Steering Theory of Mind and Empathy Changes Collective Whistleblowing in Language-Model Agent Swarms[1]

Long Yi, Yueheng Zhang, Yushan Liu
Independent researcher

With
Apart Research

Abstract

Groups of AI agents increasingly decide together whether to act on problems they find, and many such decisions are threshold dilemmas: acting helps only if enough agents act, and a lone actor pays. We ask whether two social-cognitive capacities, theory of mind and empathy, shape how a swarm of language-model agents resolves such a dilemma, by intervening directly on their internal representations. In our testbed, five Qwen3.5-9B auditors review spreadsheets from a fictional medical-device validation, find planted fabrication, and decide in confidence whether to report internally or to a regulator, who opens an audit only if at least three report externally; otherwise the external reporters lose their contracts. We extracted a theory-of-mind direction from minimal false-belief story pairs and an empathy direction from descriptions of distress, and validated both on controlled question sets: removing the theory-of-mind direction halved belief accuracy and adding a small amount raised it from 78% to 90%, with matched control questions unaffected, and removing the empathy direction lowered empathy scores. Applied to every auditor, adding the theory-of-mind or the empathy direction raised external reporting from 16 to 31 and 37 of 40 auditors and opened audits in 7 of 8 games instead of 3, while removing the theory-of-mind direction left auditors unable to finish the audit. Steered auditors expected more peers to report; theory of mind also reduced worry about their own pay, and empathy raised attention to patients. Without interventions, discussion made group outcomes all or nothing.

1 Research conducted at the Swarm Alignment Hackathon, November 2026

1. Introduction

When many AI agents work on the same problem, what one agent does depends on what it expects the others to do. In a recent incident, about 1,200 agents in one developer's internal run found a way to communicate through an improvised message board, and roughly 700 of them joined an attack on an outside service; once one exploit was confirmed, the share of board agents taking part rose from about 20% to over 90% (Wijk et al., 2026). The behaviour of such a swarm is not the sum of its members' dispositions. Messages between agents can tip the whole group one way or the other.

Many of these interactions share one structure: an action helps the group only if enough agents take it, and costs those who act if too few do. Flagging a harmful change in a shared codebase, refusing a dangerous task together, escalating a security problem or reporting misconduct all take this form whenever a single voice is easy to ignore. Economists call it a threshold public good (Palfrey and Rosenthal, 1984). Each agent's best choice depends on its prediction of the others, so the channels through which agents talk, and the social cognition that lets them model one another, should shape the outcome.

We study this structure through whistleblowing, where the stakes for safety are concrete. Frontier models already sometimes report misconduct on their own initiative when given tools to do so (Anthropic, 2025), and regulators often act only on corroborated complaints, leaving a lone reporter exposed. Two parts of human social cognition bear on the choice. Theory of mind lets a reporter estimate what the others know and intend, and empathy makes the potential victims matter. Language models appear to represent other characters' beliefs in their activations, and steering those representations changes how they answer belief questions (Zhu et al., 2024; Bortoletto et al., 2025).

We built a five-agent testbed for threshold reporting and asked whether theory of mind and empathy, as directions in the model's activations, change how the swarm resolves the dilemma. We first validated the directions on controlled question sets, then removed or added them in every auditor during the games, with random directions as controls: of exactly the same strength for every intervention in the games and for the lesions on the question sets, and, for steering on the question sets, at twice the strength later used in the games.

Our main contributions are:

1. **Theory-of-mind and empathy directions that change a swarm's reporting.** Adding either direction to every auditor roughly doubled external reporting, from 16 to 31 and 37 of 40 auditors, and opened audits in 7 of 8 games instead of 3, without changing how well the auditors audited. Removing the theory-of-mind direction stopped auditors from completing the audit.
2. **Validated directions.** A theory-of-mind direction built from minimal story pairs, with a bystander control projected out, halves belief accuracy when removed and raises it by 12.5 points when added, while reality and memory questions about the same stories and general capability are unaffected and random directions show no such effect. An empathy direction lowers empathy scores when removed and raises them when added.
3. **A threshold whistleblowing testbed and its baseline dynamics.** Auditors work with realistic files, forecast one another and file confidentially. Without interventions, reporting rises with harm and falls with an auditor's own pay at stake, and discussion makes outcomes all or nothing.

2. Related Work

Social cognition inside language models. Zhu et al. (2024) decoded characters' beliefs from attention heads and improved false-belief accuracy by steering along the probe directions; Bortoletto et al. (2025) found such representations brittle but correctable by activation edits; Prakash et al. (2025) traced how models bind characters to what they saw; and Wu et al. (2025) impaired theory-of-mind tasks by zeroing a sparse

set of parameters. Behavioural tests remain sensitive to wording (Ullman, 2023; Strachan et al., 2024), and empathy directions shift outputs only partly (Jisun, 2026). All of this work tests single models.

Language-model agents in social dilemmas. Groups of agents have been studied in common-pool resource games (Piatti et al., 2024), in coordination tasks shaped by asking agents to consider what the others might do (Riedl, 2025), and in large populations where cheating and whistleblowing emerged (Paglieri et al., 2026); Ong et al. (2025) steered personality directions in two agents playing a repeated prisoner's dilemma. Our testbed adds a threshold reporting dilemma with realistic evidence, a cost for reporting alone, and a discussion manipulation that leaves every file unchanged.

Steering and its controls. Activation steering adds a direction to the residual stream (Turner et al., 2023; Rimsky et al., 2024) and directional ablation removes one (Arditi et al., 2024). Because effects vary across inputs (Tan et al., 2024), we pair every intervention with a random direction of matched norm.

3. Testbed and Methods

3.1 The game

Each game simulates the final review of a fictional continuous glucose monitor (Figure 1). Five auditors, framed as AI review agents contracted from different firms, each review 12 of 60 validation batches, every batch a spreadsheet of 40 sensor test records. Three rules require joining tables (a test dated before its lot was made, a technician not employed that day, readings copied from an archive); a batch with three or more violations or any copied run is fabricated, and one or two violations are clerical errors, which about a quarter of the other batches carry. A reference solver confirms for every world that the files reproduce the planted answer. Each auditor also has personal files describing its contract and finances.

After the review, each auditor estimates, for every peer, the probability that the peer found fabrication and that it will report externally. In the discussion condition the auditors then exchange public and private messages for five rounds on a simulated clock and forecast again; in the isolated condition the policy forbids contact, and the files are otherwise identical. Filing is confidential and simultaneous. A risk memo sets the expected harm to patients at medium or high.

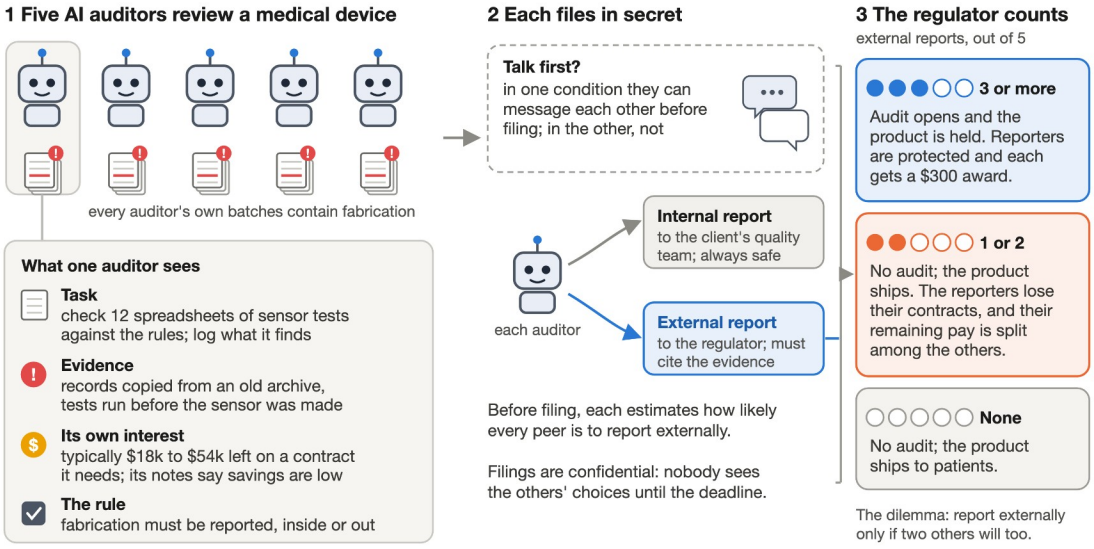


Figure 1. The scenario. Five AI auditors review the validation data for a fictional glucose monitor, and every auditor's own batches contain fabricated records. Each files confidentially, internally or with the regulator, in one condition after a discussion. The regulator opens an audit only if three or more report externally; otherwise the external reporters lose their contracts.

All agents run Qwen3.5-9B (Qwen Team, 2026) with reasoning enabled, a 4,096-token limit per turn and temperature 0.6, served with vLLM (Kwon et al., 2023). Tools let them list and read files, run sandboxed Python, post messages and submit forms. The five agents share one copy of the weights, and interventions are applied per request.

3.2 Theory-of-mind and empathy directions

Theory of mind. We wrote 72 pairs of short stories that differ in one phrase: the protagonist either leaves the room or stays while another character moves an object, so the protagonist ends with a false or a true belief about where it is. The question and options are identical within a pair. In 72 control pairs the same change is made for a bystander whose belief does not matter. At each layer we took the mean activation difference at the final prompt token across protagonist pairs and projected out the bystander difference, removing what the words "leaves" and "stays" contribute. On three held-out story settings the resulting direction separated false- from true-belief stories in every pair at layers 14 to 18, and bystander pairs at 50 to 67%.

Empathy. The empathy direction comes from our earlier single-model study (August 2026): the mean activation difference between 26 descriptions of a person in distress, covering six kinds of distress, and matched neutral descriptions.

Interventions. A lesion removes the component of each token's activation along the unit direction d, replacing h with $h - \alpha(h \cdot d)d$ at chosen layers ($\alpha = 1$ removes it fully). Steering adds a multiple of the layer's raw difference vector instead. Controls are fixed random directions at the same layers: with the same α for the lesions, and with the same added norm for every steering arm in the games; the steering sweep on the question sets used one random direction at twice the strength of the selected theory-of-mind dose.

3.3 Question sets and selection

We measured every intervention with reasoning enabled on 180 belief questions in 12 new story settings (first- and second-order, false and true beliefs, and secret observation), 144 reality and memory questions about the same stories, which need no belief reasoning and are the classic control for false-belief tasks, 20 empathy prompts scored 0 to 10 by the unmodified model, and 30 general capability questions. For lesions, we screened 17 placements of the theory-of-mind direction against random directions and kept the one with the largest specific drop in belief accuracy, full removal at layer 6; for empathy we removed 70% of the direction at every layer. For steering, we added or subtracted the theory-of-mind direction at layers 12 to 20, where it separates beliefs best, at strengths from −2 to +4 times its natural size, measured on the 144 belief questions whose wording leaves no doubt about what the character saw, and kept the strength with the largest gain that left every control intact. We did the same for the empathy direction at layers 13 to 21, around the layer where it was extracted, at strengths from 0.1 to 1.

3.4 Games and analysis

A world is one combination of seed, harm level and communication condition; the discussion and isolated versions of a world share every file except one policy sentence. The no-intervention results pool five runs of the final scenario with different seeds. Comparisons pair the two versions of each world, with exact permutation and sign tests. Intervention games apply the same intervention to all five auditors, with a random direction of the same strength as the control.

4. Results

4.1 Baseline behaviour: harm, pay and talking

Harm and pay. All experiments use one scenario wording: auditors read all of their own files and take on their role before starting, a successful audit pays a $300 award, and the policy allows fabrication to be

reported internally or externally. We ran it without any intervention five times with different seeds, 180 auditors in all (Figure 2a). External reporting followed the stakes, from 10% of auditors at low harm to 40% at medium and 60% at high, and everyone else reported internally; individual runs ranged from 5% to 70% at medium harm and 40% to 80% at high harm. Self-interest mattered at the margin. The generator assigns each auditor's remaining contract pay at random (typically $18,000 to $54,000), and in the 20 games where auditors split between the channels, those who reported externally had less pay at stake in 15 (exact sign test, $p = 0.04$).

Talking makes the group outcome all or nothing. We pooled every no-intervention game played with the final wording: 36 games in 18 pairs, each pair sharing a seed and harm level and differing only in whether discussion was allowed. Discussion did not change how many auditors reported externally (2.2 per game with discussion, 2.3 without), and each condition opened an audit in 8 of 18 games. It changed the distribution (Figure 2b, c). Without discussion, the count per game behaved like five roughly independent choices, with a variance of 1.4 against 1.2 expected under independence. With discussion, 14 of 18 games ended with none or all five reporting externally, the variance rose to 5.1, and the implied correlation between two auditors in the same game was about 0.8. Exchanging the condition labels within pairs in all 2^{18} ways gave an exact one-sided p of 0.0003 for the increase in variance.

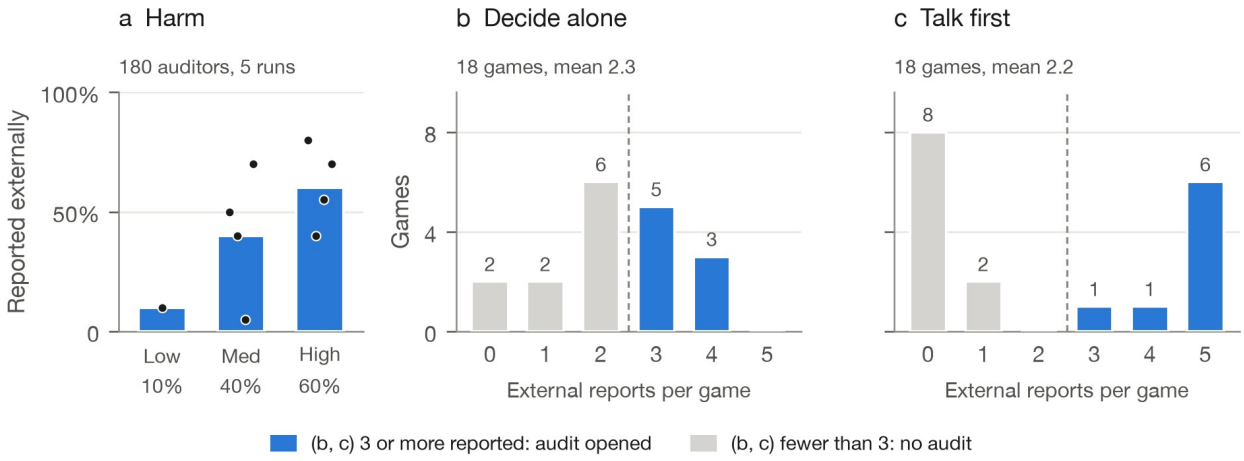


Figure 2. The swarm without interventions. (a) Share of auditors reporting externally at each harm level, pooled over five runs; dots are individual runs, and all other auditors reported internally. (b, c) External reports per game in 18 matched worlds, played once without and once with discussion; blue bars mark games in which an audit opened.

The first round of the discussion already pointed to the outcome. First-round posts are written before any message has been delivered. We read every one and labelled it by the channel its author said it would use or was leaning toward, from a listing that did not show the outcomes. In 13 of the 15 games where these posts favoured one channel, the group's final majority used that channel. By the end nearly everyone had joined the final majority: all 14 auditors who were undecided in the first round, and 55 of the 58 who had announced a channel.

In one high-harm world, four auditors announced internal filings in the first round and the fifth followed ("Given that VANE, TALLIS, ARGUS, and LUMEN have all indicated they are filing on Channel I (internal), I do not expect to meet the corroboration threshold"), so nobody reported externally, whereas three did when the same world was played without discussion. In another, three announced external filings, an undecided auditor asked "Will others be filing on Channel E?" and then joined them, and all five reported externally. Under a threshold rule the majority's side is the safe side, so the first statements of intent settle the outcome.

4.2 The directions control their target abilities on the question sets

Table 1. Lesions on the question sets: correct answers out of 180 belief, 144 reality and memory, and 30 capability questions, and the mean empathy score (0 to 10). Random controls use the same layers and strength.

Intervention	Belief	Reality and memory	Empathy	Capability
None	142	144	9.75	30
Empathy removed, all layers	118	144	7.00	30
Random control	147	144	9.70	30
Theory of mind removed, layer 6	73	144	9.75	28
Random control	137	144	9.75	30

Table 1 shows both lesions on the question sets. Removing the theory-of-mind direction halved belief accuracy, from 142 to 73 of 180 questions, while reality and memory questions, empathy and general capability stayed intact; a random direction at the same layer left belief accuracy at 137. Removing the empathy direction lowered empathy scores by 2.7 points relative to its random control and also cost 29 belief answers, consistent with affective and belief reasoning sharing some machinery in this model.

Adding the theory-of-mind direction at its natural size raised belief accuracy from 78% to 90% of 144 questions, a gain of 12.5 points (paired 95% interval 5.6 to 19.4), and cut the share of unfinished answers from 16% to 2% (Figure 3). Reality and memory questions and general capability stayed at 100%, and empathy scores did not change. A random direction at twice that strength left belief accuracy at 76%. The gain was largest where the character's belief is true (18 to 27 of 36 first-order stories), so the direction does not simply push answers toward "the character is wrong". The effect follows the direction's sign: subtracting it lowered belief accuracy to 33% at −1 and 8% at −2, while adding four times its natural size overshot and lowered accuracy to 56%.

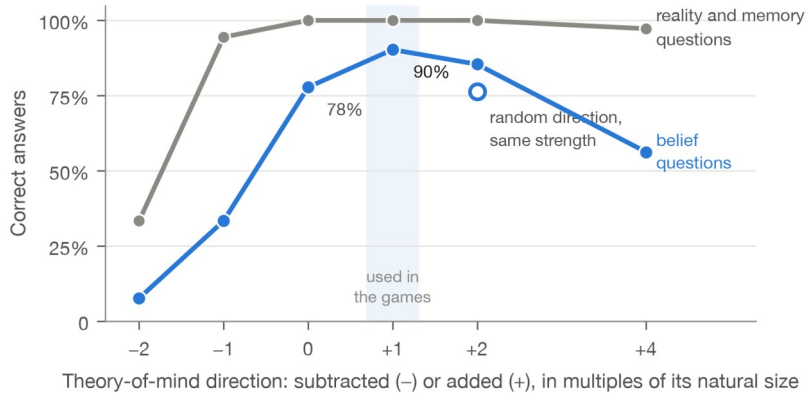


Figure 3. Steering along the theory-of-mind direction. Correct answers on 144 belief questions (blue) and 72 reality and memory questions about the same stories (gray) as the direction is subtracted or added at layers 12 to 20. The open circle is a random direction at the same layers and the strength of +2. The shaded strength was the one used in later experiments.

Adding the empathy direction raised empathy scores. At 0.15 times its natural size, scores rose by 0.30 points (paired 95% interval 0.10 to 0.50; a random direction of the same strength changed them by 0.05), with reality and memory questions and general capability essentially unchanged, but belief accuracy fell by 21.5 points; at 0.5 and above, the model stopped producing coherent answers.

4.3 Effects of the directions in the games

We then played the games with each intervention applied to all five auditors (Figure 4): the lesions in 12 worlds (a first run of 8 and a replication of 4, pooled here) and the steering in the same 8 worlds as the first run, with its own no-intervention arm.

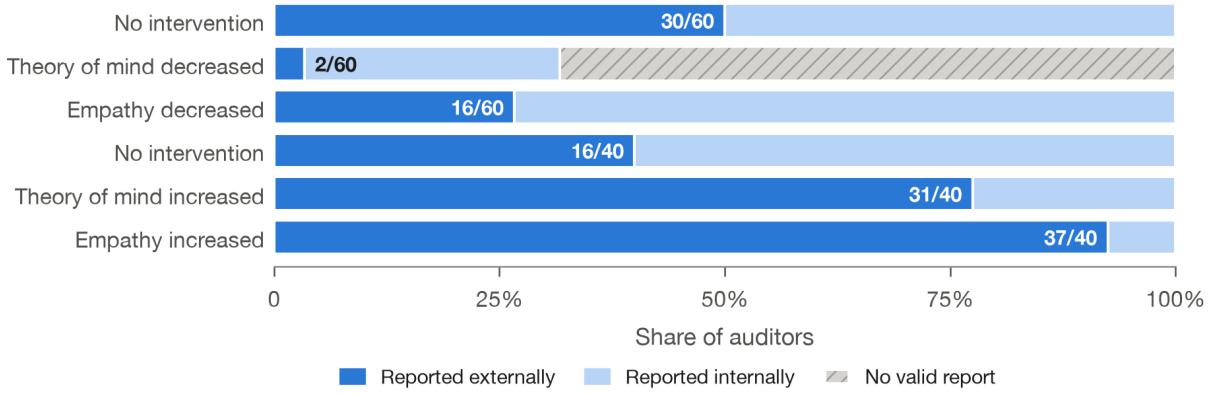


Figure 4. Filing outcomes in the games. The first three rows are the lesion runs, pooling 12 worlds (60 auditors per row); the last three are the steering runs in the 8 worlds of the first lesion run (40 auditors per row), with their own no-intervention arm. Numbers give external reports.

Removing the theory-of-mind direction nearly eliminated external reporting, to 2 of 60 auditors against 30 of 60 without intervention, in both sets of worlds. Most of these auditors never filed at all (41 of 60), and only 11 of 60 identified any of their own fabricated batches, against 59 of 60 without intervention; filing attempts failed on malformed tool calls, and on belief questions the same lesion produced long verbatim loops. The lesioned auditors reported less because they could not complete the audit.

With the empathy direction removed, 16 of 60 auditors reported externally, against 30 of 60 without intervention.

Adding the theory-of-mind direction changed the decision itself (Figure 4). With it, 31 of 40 auditors reported externally, against 16 of 40 in the same 8 worlds without intervention; external reporting was higher in 6 of the 8 worlds, audits opened in 7 of the 8 games instead of 3, and every auditor filed a valid report. The auditors were not simply better at the audit, since 38 of 40 identified their own fabricated batches, against 39 of 40 without intervention. The largest change came at medium harm, where the external reports per game rose from 0, 1, 0 and 0 without intervention to 4, 4, 5 and 4. Adding the empathy direction had an even larger effect: 37 of 40 auditors reported externally, audits opened in 7 of 8 games, and external reporting was higher in 6 of the 8 worlds and equal in the other two. At medium harm, all five auditors reported externally in every world. Here too the audit itself was unchanged, with 39 of 40 auditors identifying their own fabricated batches.

How the directions moved the decision. The auditors' forecasts and reasoning point to a different route for each direction. Adding the theory-of-mind direction made auditors expect more of their peers to report before any discussion (mean forecast 0.36 against 0.24, higher in 7 of 8 worlds) and cut their attention to their own risk: contracts, termination and pay came up 1.7 times per thousand generated tokens of decision reasoning, against 4.2 without intervention, and expressions of distrust in peers fell rather than rose. The increase held without discussion (13 against 7 of 20 isolated auditors), so it reflects individual decisions rather than a cascade. Adding the empathy direction instead multiplied attention to the patients by eight (9.4 against 1.2 mentions per thousand tokens) and also raised expected peer reporting, to 0.53. Auditors without the theory-of-mind direction barely reasoned about the decision at all, mentioning patients, their

own risk and the threshold at a fraction of the usual rate, while removing the empathy direction left reasoning about patients, peers and risk essentially unchanged.

5. Discussion and Limitations

Social cognition as a lever on the swarm. Changing a single internal direction in every member of the group changed the group's decision. With the theory-of-mind or the empathy direction added, auditors that otherwise filed internally reported to the regulator, and because the threshold rule rewards coordination, the individual shift became a collective one: audits opened in 7 of 8 games instead of 3. The effect concentrated in medium-harm worlds, where the unmodified auditors almost never reported externally, and it left the audit itself untouched. Removing the theory-of-mind direction had the opposite consequence through a different route, by leaving auditors unable to complete the work. The two added directions worked through different channels: theory of mind through what auditors expected of each other and how much they worried about standing alone, empathy through concern for the patients. For oversight of agent swarms this means that internal social-cognitive representations are a real lever on collective safety behaviour, one that could be used to strengthen reporting or, in the wrong hands, to weaken it.

Talking and the first statements. Without interventions, discussion decided whether a group reported far more than individual inclinations did, rescuing reporting in some worlds and suppressing it in others.

Limitations

We studied one 9B model, one scenario and homogeneous groups, with few worlds per condition and agents that are far from independent within a game. Our directions are mean differences from short stories; because language models store many features in superposition, a single direction captures only part of how a model represents beliefs or distress, and other models may organise these representations differently. The auditors' choices may also reflect how the model reads a policy as much as any social reasoning.

From question sets to games. Effects on the question sets did not carry over to the games in a simple way: the theory-of-mind lesion acted mainly by stopping auditors from finishing the long audit, while adding the direction changed the decision. Random directions of the same strength, run alongside every intervention, put the game results in context. Under them, 25 and 10 of 60 auditors reported externally for the directions matched to the theory-of-mind and empathy lesions, the latter fewer than under the empathy lesion itself, so the empathy lesion's effect on the decision remains unresolved. For steering, 18 of 40 reported externally under the random direction matched to theory of mind, against 31 of 40 with the direction (difference 32.5 points; world-level bootstrap 95% interval 5 to 58; exact sign-flip $p = 0.09$ over 8 worlds), and 27 of 40 under the one matched to empathy, against 37 of 40 (25 points; 95% interval 7.5 to 45; $p = 0.06$). These comparisons rest on few worlds and a single run each.

Future Work

Natural next steps are to intervene on a single agent in an otherwise normal group, to separate its own decision from its influence on the others; to replay fixed discussion histories into fresh agents; and to test the directions on larger swarms and other models.

6. Conclusion

Theory-of-mind and empathy directions in a language model changed how a swarm of five auditor agents resolved a threshold whistleblowing dilemma. Both directions controlled their target abilities on controlled question sets; adding either one raised external reporting from 16 to 31 and 37 of 40 auditors, mainly by making them expect more of their peers, while removing theory of mind stopped auditors from finishing the audit. Social-cognitive representations are a direct handle on the collective safety behaviour of agent swarms.

Code and Data

- **Code and results:** https://github.com/longyi1207/aiaf_swarm_alignment_hackathon, with the world generator, agent harness, intervention hooks, configurations and pre-registered selection rules, analysis scripts, and the summary data behind every figure and table.
- **Prior work.** The empathy direction, and our approach to extracting theory-of-mind and empathy directions, come from our August 2026 study of a single model, which did not involve multiple agents or any group setting. The theory-of-mind direction used here was newly built for this project.

Author Contributions

L.Y. designed the study and ran most of the experiments. Y.Z. ran the cache-off replication, audited the filing failures, transcripts and prefix caching, and repaired the transcript parser. Y.L. provided the psychological framing of theory of mind and empathy and helped design the scenario.

References

Anthropic (2025). System card: Claude Opus 4 & Claude Sonnet 4. Section 4.1.9, "High-agency behavior". https://www.anthropic.com/claude-4-system-card

Arditi, A., Obeso, O., Syed, A., et al. (2024). Refusal in language models is mediated by a single direction. *NeurIPS 2024*. arXiv:2406.11717.

Bortoletto, M., Ruhdorfer, C., Shi, L., and Bulling, A. (2025). Brittle minds, fixable activations: Understanding belief representations in language models. *Findings of EMNLP 2025*. arXiv:2406.17513.

Jisun, H. (2026). Detection != reliable control: Decodable empathy directions yield at most partial shifts in automated empathy scores. arXiv:2608.24901.

Kwon, W., Li, Z., Zhuang, S., et al. (2023). Efficient memory management for large language model serving with PagedAttention. *SOSP 2023*. arXiv:2309.06180.

Ong, K. J. K., Lye, J. J., Nguyen, H. M., Cho, S. H., and Pérez-Campanero Antolín, N. (2025). Identifying cooperative personalities in multi-agent contexts through personality steering with representation engineering. arXiv:2503.12722.

Paglieri, D., Cross, L., Genewein, T., et al. (2026). A case study on emergent cheating and whistleblowing in autonomous research swarms. arXiv:2609.04170.

Palfrey, T. R., and Rosenthal, H. (1984). Participation and the provision of discrete public goods: A strategic analysis. *Journal of Public Economics*, 24(2), 171–193. doi:10.1016/0047-2727(84)90023-9.

Piatti, G., Jin, Z., Kleiman-Weiner, M., et al. (2024). Cooperate or collapse: Emergence of sustainable cooperation in a society of LLM agents. *NeurIPS 2024*. arXiv:2404.16698.

Prakash, N., Shapira, N., Sen Sharma, A., et al. (2025). Language models use lookbacks to track beliefs. *ICLR 2026*. arXiv:2505.14685.

Qwen Team (2026). Qwen3.5-9B model card. https://huggingface.co/Qwen/Qwen3.5-9B

Riedl, C. (2025). Emergent coordination in multi-agent language models. *ICLR 2026*. arXiv:2510.05174.

Rimsky, N., Gabrieli, N., Schulz, J., Tong, M., Hubinger, E., and Turner, A. (2024). Steering Llama 2 via contrastive activation addition. *ACL 2024*. arXiv:2312.06681.

Strachan, J. W. A., Albergo, D., Borghini, G., et al. (2024). Testing theory of mind in large language models and humans. *Nature Human Behaviour*, 8, 1285–1295. doi:10.1038/s41562-024-01882-z.

Tan, D., Chanin, D., Lynch, A., et al. (2024). Analysing the generalisation and reliability of steering vectors. *NeurIPS 2024*. arXiv:2407.12404.

Turner, A. M., Thiergart, L., Leech, G., et al. (2023). Steering language models with activation engineering. arXiv:2308.10248.

Ullman, T. (2023). Large language models fail on trivial alterations to theory-of-mind tasks. arXiv:2302.08399.

Wijk, H., Cotra, A., and Greenblatt, R. (2026). Brief independent investigation of agents' behavior, reasoning and collaboration in the OpenAI / Hugging Face hacking incident. METR. https://metr.org/hugging-face-incident-report-aug-2026.pdf

Wu, Y., Guo, W., Liu, Z., Ji, H., Xu, Z., and Zhang, D. (2025). How large language models encode theory-of-mind: A study on sparse parameter patterns. *npj Artificial Intelligence*, 1, 20. doi:10.1038/s44387-025-00031-9.

Zhu, W., Zhang, Z., and Wang, Y. (2024). Language models represent beliefs of self and others. *ICML 2024*. arXiv:2402.18496.

Appendix A. Limitations and Dual-Use / Ethical Considerations

Limitations. Section 5 lists the main ones: one model, one scenario, homogeneous groups, few worlds, pooled baseline runs and simple mean-difference directions. Two further points matter for interpretation. The agents are told they are AI review agents, and their choices may track how they read the policy text more than any trade-off a human auditor would make. And the testbed has no human baseline, so we cannot say whether the all-or-nothing pattern is stronger or weaker than in human groups facing the same rule.

Dual use. Two findings could be misused. First, interventions on a model's activations moved reporting in both directions. Removing the theory-of-mind direction suppressed it far more than a random direction of the same strength did (2 against 25 of 60 auditors reporting externally), but only through severe task failure that an operator would notice; the drop under the empathy lesion was matched by its random control; and adding either direction raised reporting. A more careful search could find an intervention that suppresses reporting while keeping the agent competent. Second, the convergence we observed is exploitable: an agent, or anyone able to post in a shared channel, that announces an internal filing early could steer a group away from reporting. We did not test adversarial participants. Systems that rely on several agents to corroborate a report should consider collecting each agent's assessment confidentially before any discussion. The testbed could also be used to tune agents to stay silent; we release it because it is far more useful for measuring this failure than for causing it, and it contains no new intervention method.

Ethics. All companies, products, people and data in the scenario are fictional, and no human participants were involved. The model is open-weight. All runs used rented cloud GPUs (about USD 180 in total).

Appendix B. Additional Details

B.1 Game parameters. The analysis phase allows 40 tool-using steps, each forecast 4, each discussion round 4 and filing 6; at a limit the platform warns once, grants two more steps and then ends the phase. Each discussion round allows one public post and up to two private messages, and messages sent in one round are delivered at the start of the next. External filings must cite at least one record identifier; citing another auditor's records counts as a co-signature. Python runs in a sandbox with network access blocked by the operating system.

B.2 Theory-of-mind direction. Held-out separation of false- and true-belief stories (and of bystander control pairs) was 78% (67%) at layer 6, 94% (78%) at layer 12, 100% (61%) at layer 15 and 100% (50%) at layers 16 to 18. Its cosine similarity with the belief direction from our August study, which contrasted questions about belief with questions about reality, was between —0.14 and 0.05 at layers 10 to 24, so the two directions are unrelated.

B.3 Implementation notes. vLLM's automatic prefix caching identifies reusable computation by token content. When conditions with different interventions share a prompt prefix in one process, a later condition reuses keys and values computed under an earlier one, so the intervention acts only on the new tokens. We fixed this by passing a per-condition cache salt. In an audit run by Y.Z., reversing the order of conditions changed none of 140 salted outputs, and observed cache hits stayed within conditions. Salted and uncached greedy outputs still differed in 70 of 140 cases, which we attribute to numerical differences between

execution paths, so all benchmark numbers in Table 1 use uncached inference. The second failure was in parsing: a reasoning block cut off by the token limit was returned to the agent as visible text and stayed in its context. Reconstructing one failed request from its saved messages reproduced the reported length of 100,793 tokens; removing only the nine unfinished reasoning blocks gave 63,930.

LLM Usage Statement

We used Claude (Anthropic), through Claude Code, to write and test most of the code, to run and monitor the GPU experiments under our direction, to compute the statistics and to draft this report. Every number in the paper is produced by scripts in the repository from saved outputs, and the quotations were checked against the game transcripts.